

Completed

Norah Jones

2026-02-09

Table of contents

Preface	3
1 Project_One	4
2 Day3 Prac Questions 4 and 5	8
3 Practical Day Four	11

Preface

This is a Quarto book.

To learn more about Quarto books visit <https://quarto.org/docs/books>.

1 + 1

[1] 2

1 Project_One

QUESTION ONE

```
down <- nrow(airquality)
across <- ncol(airquality)
sum <- 0

print('The following rows have missing information: ')
```

```
[1] "The following rows have missing information: "
```

```
for(i in 1:down){
  for(j in 1:across){
    if(is.na(airquality[i,j])){
      sum <- sum + 1
      print(i)
      break}
  }
}
```

```
[1] 5
[1] 6
[1] 10
[1] 11
[1] 25
[1] 26
[1] 27
[1] 32
[1] 33
[1] 34
[1] 35
[1] 36
[1] 37
[1] 39
```

```
[1] 42  
[1] 43  
[1] 45  
[1] 46  
[1] 52  
[1] 53  
[1] 54  
[1] 55  
[1] 56  
[1] 57  
[1] 58  
[1] 59  
[1] 60  
[1] 61  
[1] 65  
[1] 72  
[1] 75  
[1] 83  
[1] 84  
[1] 96  
[1] 97  
[1] 98  
[1] 102  
[1] 103  
[1] 107  
[1] 115  
[1] 119  
[1] 150
```

```
print(paste0('In total there are ', sum, ' rows with missing information.' ))
```

```
[1] "In total there are 42 rows with missing information."
```

QUESTION TWO

```
my_table <- data.frame(  
  Column = c('Temperature', 'Ozone'),  
  Mean = c(mean(airquality[,4]), mean(airquality[,1], na.rm=TRUE)),  
  SD = c(sd(airquality[,4]), sd(airquality[,1], na.rm=TRUE)),  
  Min = c(min(airquality[,4]), min(airquality[,1], na.rm=TRUE)),  
  Max = c(max(airquality[,4]), max(airquality[,1], na.rm=TRUE))
```

```
)  
  
print(my_table)
```

	Column	Mean	SD	Min	Max
1	Temperature	77.88235	9.46527	56	97
2	Ozone	42.12931	32.98788	1	168

QUESTION THREE

```
data(cars)  
  
Y <- cars$dist  
X <- cbind(1, cars$speed)  
  
my_funct <- function(design, response) {  
  
  a <- t(design)%*%design  
  b <- solve(a)  
  c <- t(design)%*%response  
  d <- b%*%c  
  
  return(d)  
}  
  
print(my_funct(X, Y))
```

```
[,1]  
[1,] -17.579095  
[2,] 3.932409
```

QUESTION FOUR

```
model <- lm(cars$dist~cars$speed, data=cars )  
summary(model)
```

Call:
lm(formula = cars\$dist ~ cars\$speed, data = cars)

Residuals:

Min	1Q	Median	3Q	Max
-29.069	-9.525	-2.272	9.215	43.201

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-17.5791	6.7584	-2.601	0.0123 *
cars\$speed	3.9324	0.4155	9.464	1.49e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 15.38 on 48 degrees of freedom

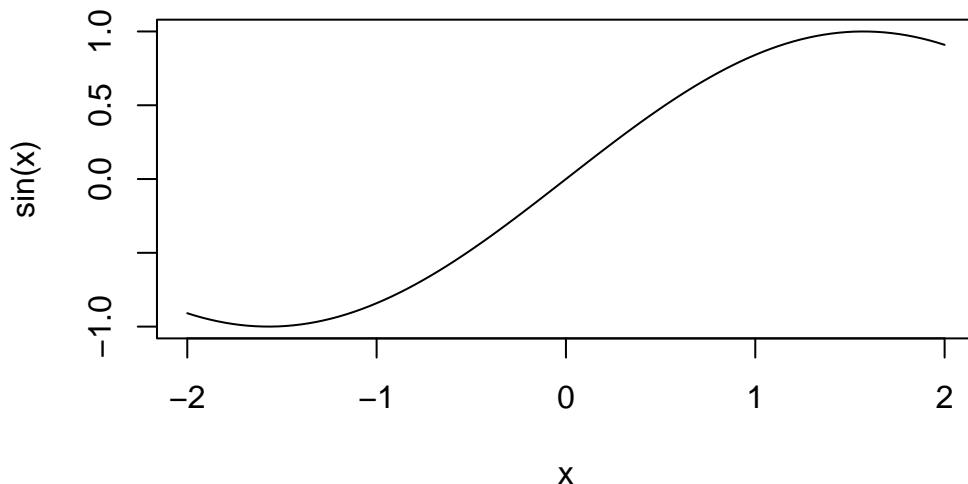
Multiple R-squared: 0.6511, Adjusted R-squared: 0.6438

F-statistic: 89.57 on 1 and 48 DF, p-value: 1.49e-12

2 Day3 Prac Questions 4 and 5

QUESTION FOUR

```
curve(sin(x), from = -2, to = 2)
```



QUESTION FIVE

```
randoms <- rt(1000, 1)

#Creating a manual QQ-plot with 95% confidence interval

sorted_randoms <- sort(randoms)

n <- length(randoms)
i <- 1:n
```

```

# Calculate plotting positions (probabilities)
p <- (i - 3/8) / (n + 1/4)

# Generate normal quantiles by simulation (since we can't use qnorm)
large_normal_sample <- rnorm(1000000) # Large normal reference
theoretical_quantiles <- quantile(large_normal_sample, probs = p)

# Calculate 95% probability envelopes by simulation
n_sim <- 1000 # Number of simulations
envelope_matrix <- matrix(NA, nrow = n_sim, ncol = n)

# Simulate many normal samples of size n
for (j in 1:n_sim) {
  sim_sample <- rnorm(n)
  sim_sorted <- sort(sim_sample)
  envelope_matrix[j, ] <- sim_sorted
}

# Calculate 2.5% and 97.5% percentiles at each position
lower_envelope <- apply(envelope_matrix, 2, quantile, probs = 0.025)
upper_envelope <- apply(envelope_matrix, 2, quantile, probs = 0.975)

# Create the QQ-plot
par(mfrow = c(1, 2)) # Split plot window

# Manual QQ-plot
plot(theoretical_quantiles, sorted_randoms,
      xlab = "Theoretical Normal Quantiles",
      ylab = "Sample Quantiles",
      main = "Manual QQ-plot with 95% Envelopes",
      pch = 19, cex = 0.6,
      ylim = c(-5, 5), xlim = c(-4, 4),
      col = rgb(0, 0, 0, 0.7))

# Add reference line (y = x)
abline(a = 0, b = 1, col = "red", lwd = 2)

# Add 95% probability bands
lines(theoretical_quantiles, lower_envelope,
      col = "blue", lty = 2, lwd = 2)
lines(theoretical_quantiles, upper_envelope,
      col = "blue", lty = 2, lwd = 2)

```

```

# Add legend
legend("topleft",
       legend = c("Data", "Reference Line", "95% Envelope"),
       col = c("black", "red", "blue"),
       pch = c(19, NA, NA),
       lty = c(NA, 1, 2),
       lwd = c(NA, 2, 2))

# Check with car::qqPlot for comparison
library(car)

```

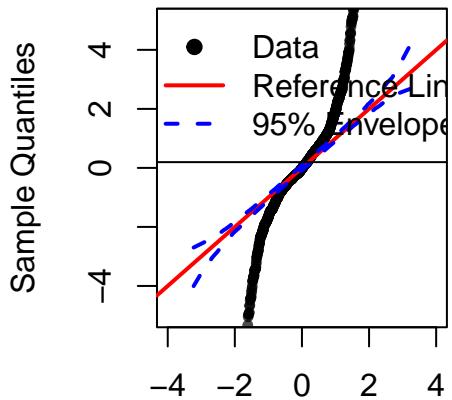
Loading required package: carData

```

qqPlot(randoms, envelope = 0.95, ylim = c(-5, 5),
       main = "car::qqPlot for Comparison",
       xlab = "Normal Quantiles",
       ylab = "Sample Quantiles")

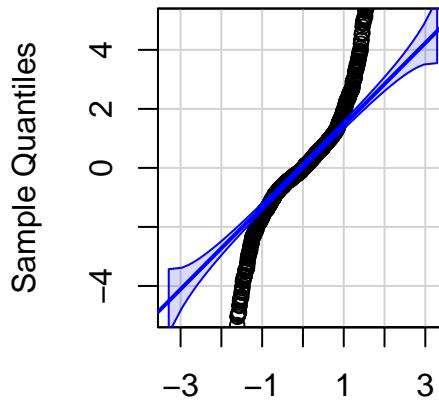
```

Manual QQ-plot with 95% Envelope



Theoretical Normal Quantiles

car::qqPlot for Comparison



Normal Quantiles

[1] 657 94

3 Practical Day Four

3.1

QUESTION ONE

```
library(tidyverse)
```

Warning: package 'tidyverse' was built under R version 4.3.3

Warning: package 'tidyr' was built under R version 4.3.3

Warning: package 'purrr' was built under R version 4.3.3

Warning: package 'dplyr' was built under R version 4.3.3

Warning: package 'forcats' was built under R version 4.3.3

Warning: package 'lubridate' was built under R version 4.3.3

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
v forcats   1.0.0     v stringr   1.6.0
v ggplot2   4.0.2     v tibble    3.2.1
v lubridate 1.9.4     v tidyr    1.3.1
v purrr    1.0.4
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(nycflights13)
```

Warning: package 'nycflights13' was built under R version 4.3.3

```
#The line below renames the first 10 columns of flights and prints the flights data set with
flights |> rename(Year = year, Month = month, Day = day, Dep_Time = dep_time, Sched_Dep_Time

# A tibble: 336,776 x 19
  Year Month   Day Dep_Time Sched_Dep_Time Dep_Delay Arr_Time Sched_Arr_Time
  <int> <int> <int>     <int>           <int>     <dbl>     <int>           <int>
1 2013     1     2     2130         2130       0         1           18
2 2013     1    11     2157         2000      117        1          2208
3 2013     1    11     2253         2249       4         1          2357
4 2013     1    14     2122         2130      -8         1            2
5 2013     1    14     2246         2250      -4         1            7
6 2013     1    15     2304         2245      19         1          2357
7 2013     1    16     2018         2025      -7         1          2329
8 2013     1    16     2303         2245      18         1          2357
9 2013     1    19     2107         2110      -3         1          2355
10 2013    1    22     2246         2249      -3         1          2357
# i 336,766 more rows
# i 11 more variables: Arr_Delay <dbl>, Carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

QUESTION TWO

```
#Creates flight1 tibble that contains flights that occurred in month 1 ONLY
flight1 <- flights |> filter(month == 1)

#Gets the average distance of each carrier
carrier_dist_vec_mean <- flight1 |> group_by(carrier) |> summarise(MEAN = mean(distance, na.rm = TRUE))

#Gets the standard deviation of the distance for each carrier
carrier_dist_vec_sd <- flight1 |> group_by(carrier) |> summarise(SD = sd(distance, na.rm = TRUE))

#Creates one tibble with both the mean and sd distance for each carrier
dist_tbl <- carrier_dist_vec_mean |> left_join(carrier_dist_vec_sd)
```

```
Joining with `by = join_by(carrier)`
```

```
#Arranges the carrier rows from that with the least to the highest mean distance  
dist_tbl |> arrange(MEAN)
```

```
# A tibble: 16 x 3  
  carrier    MEAN     SD  
  <chr>    <dbl> <dbl>  
1 YV        229     0  
2 9E       476.   334.  
3 EV       522.   294.  
4 US       536.   553.  
5 MQ       566.   223.  
6 FL       691.   142.  
7 00       733     NA  
8 WN       942.   496.  
9 B6      1062.   681.  
10 DL     1220.   644.  
11 AA     1350.   626.  
12 UA     1462.   778.  
13 F9     1620     0  
14 AS     2402     0  
15 VX     2495.   98.2  
16 HA     4983     0
```

QUESTION THREE

```
##### NA #####
```

```
carrier00 <- flight1 |> filter(carrier == '00') |> select(carrier, distance)  
carrier00
```

```
# A tibble: 1 x 2  
  carrier distance  
  <chr>    <dbl>  
1 00        733
```

```
# Explanation of NA: Carrier 00 has NA as the standard deviation of the distance, that is be
```

```

#####
# ZERO #####
carrierYV <- flight1 |> filter(carrier == 'YV') |> select(carrier, distance)

carrierYV

# A tibble: 46 x 2
  carrier distance
  <chr>     <dbl>
1 YV         229
2 YV         229
3 YV         229
4 YV         229
5 YV         229
6 YV         229
7 YV         229
8 YV         229
9 YV         229
10 YV        229
# i 36 more rows

# Explanaiton of ZERO: The standard deviation of the distance for the YV carrier is zero because all flights are 229 miles

```

QUESTION FOUR

```

#Getting the Average Departure delay for each carrier in each month in the long format
switched <- flights |> group_by(carrier, month) |> summarise(Av_Dep_Delay = mean(dep_delay, na.rm = TRUE))

`summarise()` has grouped output by 'carrier'. You can override using the
`.groups` argument.

```

```

#Converting into wide data
switched2 <- switched |> pivot_wider(names_from = carrier, values_from = Av_Dep_Delay)

```

QUESTION FIVE

```

#Getting number of all flights with a delayed departure time and arrived with NO delay
meet_Condition <- flights |> filter(dep_delay > 0) |> filter(arr_delay <= 0) |> count()

#Getting number of total flights

```

```

total_Flights <- flights |> count()

#Calculating the proportion
prop_Delayed <- meet_Condition/total_Flights

#Renaming the column
prop_Delayed <- prop_Delayed |> rename(Prop_Delayed = n)

```

QUESTION SIX

```

routes <- flights |> group_by(origin, dest) |> summarize(n_airlines = n_distinct(carrier)) |>

`summarise()` has grouped output by 'origin'. You can override using the
`.groups` argument.

routes

# A tibble: 128 x 3
# Groups:   origin [3]
  origin dest  n_airlines
  <chr>  <chr>     <int>
1 EWR    ATL      4
2 EWR    AUS      2
3 EWR    BDL      2
4 EWR    BNA      2
5 EWR    BOS      3
6 EWR    BWI      2
7 EWR    CHS      2
8 EWR    CLE      2
9 EWR    CLT      3
10 EWR   CVG      2
# i 118 more rows

#Joining the routes with at least two airlines with rest of the flights data
joint <- routes |> left_join(flights, by = join_by(origin, dest))

#Grouping to get each route+carrier combo then calc it's average arrival delay
joint2 <- joint |> group_by(origin, dest, carrier) |> summarise(Av_Arr_Delay = mean(arr_delay))

`summarise()` has grouped output by 'origin', 'dest'. You can override using
the `.groups` argument.

```

```
joint2
```

```
# A tibble: 343 x 4
# Groups:   origin, dest [128]
  origin dest carrier Av_Arr_Delay
  <chr>  <chr>  <chr>      <dbl>
1 EWR    ATL    9E        -6.25
2 EWR    ATL    DL        10.00
3 EWR    ATL    EV        19.5
4 EWR    ATL    UA        10.5
5 EWR    AUS    UA        4.28
6 EWR    AUS    WN       -11.2
7 EWR    BDL    EV        6.78
8 EWR    BDL    UA        22.6
9 EWR    BNA    EV        17.7
10 EWR   BNA    WN       -2.13
# i 333 more rows
```

```
# Airline with the best average arrival delay
```

```
best <- joint2 |> group_by(origin, dest) |> summarise(best_av_arr_delay = min(Av_Arr_Delay))
```

```
`summarise()` has grouped output by 'origin'. You can override using the
`.groups` argument.
```

```
best2 <- best |> left_join(joint2, join_by(best_av_arr_delay == Av_Arr_Delay)) |> select(!or
```

```
Warning in left_join(best, joint2, join_by(best_av_arr_delay == Av_Arr_Delay)): Detected an
  to-many relationship between `x` and `y`.
  i Row 17 of `x` matches multiple rows in `y`.
  i Row 41 of `y` matches multiple rows in `x`.
  i If a many-to-many relationship is expected, set `relationship =
    "many-to-many"` to silence this warning.
```

```
best2
```

```
# A tibble: 134 x 4
  origin.x dest.x best_av_arr_delay carrier
  <chr>     <chr>      <dbl>  <chr>
```

```

1 EWR      ATL          -6.25  9E
2 EWR      AUS          -11.2   WN
3 EWR      BDL          6.78   EV
4 EWR      BNA          -2.13   WN
5 EWR      BOS          -4.01   EV
6 EWR      BWI          5.95   WN
7 EWR      CHS          -14     UA
8 EWR      CLE          -3.71   EV
9 EWR      CLT          0.920  US
10 EWR     CVG          1.40   9E
# i 124 more rows

```

```
#Airline with the worst arrival delay
```

```
worst <- joint2 |> group_by(origin, dest) |> summarise(worst_av_arr_delay = max(Av_Arr_Delay))
```

`summarise()` has grouped output by 'origin'. You can override using the `.`groups` argument.

```
worst2 <- worst |> left_join(joint2, join_by(worst_av_arr_delay == Av_Arr_Delay)) |> select(
```

Warning in left_join(worst, joint2, join_by(worst_av_arr_delay == Av_Arr_Delay)): Detected an
to-many relationship between `x` and `y`.
i Row 75 of `x` matches multiple rows in `y`.
i Row 194 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected, set `relationship =
"many-to-many"` to silence this warning.

```
worst2
```

```
# A tibble: 130 x 4
  origin.x dest.x worst_av_arr_delay carrier
  <chr>    <chr>        <dbl> <chr>
1 EWR      ATL          19.5  EV
2 EWR      AUS          4.28  UA
3 EWR      BDL          22.6  UA
4 EWR      BNA          17.7  EV
5 EWR      BOS          6.87  B6
6 EWR      BWI          20.1  EV
7 EWR      CHS          16.2  EV
```

```
8 EWR      CLE          5.97 UA
9 EWR      CLT          20.5  EV
10 EWR     CVG          21.2  EV
# i 120 more rows
```

```
#Combining
```

```
results <- best2 |> left_join(worst2, join_by(origin.x, dest.x))
```

```
Warning in left_join(best2, worst2, join_by(origin.x, dest.x)): Detected an unexpected many-to-many relationship between `x` and `y`.
i Row 77 of `x` matches multiple rows in `y`.
i Row 17 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected, set `relationship =
"many-to-many"` to silence this warning.
```

```
results
```

```
# A tibble: 136 x 6
  origin.x dest.x best_av_arr_delay carrier.x worst_av_arr_delay carrier.y
  <chr>     <chr>        <dbl> <chr>        <dbl> <chr>
1 EWR       ATL        -6.25  9E           19.5  EV
2 EWR       AUS       -11.2   WN           4.28  UA
3 EWR       BDL        6.78   EV           22.6  UA
4 EWR       BNA       -2.13   WN           17.7  EV
5 EWR       BOS        -4.01  EV           6.87  B6
6 EWR       BWI        5.95   WN           20.1  EV
7 EWR       CHS       -14     UA           16.2  EV
8 EWR       CLE       -3.71   EV           5.97  UA
9 EWR       CLT        0.920  US           20.5  EV
10 EWR      CVG        1.40   9E           21.2  EV
# i 126 more rows
```

```
results2 <- results |> mutate(diff = sqrt((best_av_arr_delay - worst_av_arr_delay)**2))
```

```
results3 <- results2 |> summarise(greatest_diff = max(diff, na.rm = TRUE)) |> left_join(resu
```

```
results3
```

```
# A tibble: 1 x 7
```

```

greatest_diff origin.x dest.x best_av_arr_delay carrier.x worst_av_arr_delay
<dbl> <chr>      <chr>          <dbl> <chr>          <dbl>
1       127. JFK      ATL           1.40 9E        128
# i 1 more variable: carrier.y <chr>

#Trying to find the reason

reason <- results |> mutate(diff = sqrt((best_av_arr_delay - worst_av_arr_delay)**2))

reason2 <- reason |> summarise(terrible = max(worst_av_arr_delay, na.rm = TRUE))

reason2

# A tibble: 1 x 1
  terrible
  <dbl>
1 128

# The above code shows that the worst average arrival delay was 128, it so happens that this

```

QUESTION SEVEN

```

given_data <- structure(list(id = c("id_1", "id_2", "id_3", "id_4", "id_5",
"id_6", "id_7", "id_8", "id_9", "id_10", "id_11", "id_12", "id_13",
"id_14", "id_15", "id_16", "id_17", "id_18", "id_19", "id_20",
"id_21", "id_22", "id_23", "id_24", "id_25", "id_26", "id_27",
"id_28", "id_29", "id_30", "id_31", "id_32", "id_33", "id_34",
"id_35", "id_36", "id_37", "id_38", "id_39", "id_40", "id_41",
"id_42", "id_43", "id_44", "id_45", "id_46", "id_47", "id_48",
"id_49", "id_50"), age = c(50L, 34L, 70L, 33L, 22L, 61L, 69L,
73L, 62L, 56L, 71L, 33L, 73L, 44L, 45L, 46L, 24L, 70L, 46L, 76L,
47L, 76L, 28L, 48L, 54L, 27L, 45L, 26L, 61L, 28L, 38L, 55L, 33L,
36L, 62L, 58L, 72L, 31L, 34L, 51L, 61L, 64L, 26L, 28L, 60L, 29L,
42L, 46L, 79L, 72L), gender = c("male", "male", "male", "female",
"female", "male", "female", "male", "male", "female", "female",
"male", "male", "female", "male", "male", "male", "female",
"male", "male", "male", "female", "female", "female", "male",
"female", "female", "male", "female", "female", "female",
"male", "male", "female", "male", "female", "female", "male",
"female", "female", "male", "male", "female", "male", "male",
"male", "female"), height = c(174.4, 197.7, 174.1, 194.5, NA,

```



```

"occasional", "occasional", "none", "none", "none", "none",
"occasional", "regular", "regular", "none", "regular", "occasional",
"occasional", "none", "occasional", "regular"), income = c(84820L,
81547L, 22588L, 72490L, 74533L, 25338L, 41469L, 57315L, 63629L,
88662L, 62615L, 56261L, 58499L, 82232L, 77584L, 77275L, 38468L,
54510L, 91326L, 78611L, 31402L, 29586L, 21441L, 58269L, 84173L,
88295L, 37940L, 43750L, 69750L, 92356L, 82518L, 91455L, 68866L,
51178L, 68275L, 27689L, 35418L, 81318L, 62405L, 86851L, 25654L,
47553L, 74474L, 51409L, 22607L, 55360L, 96351L, 21516L, 41927L,
55810L), education = c("master", "bachelor", "PhD", "master",
"bachelor", "highschool", "PhD", "highschool", "PhD", "PhD",
"bachelor", "highschool", "master", "bachelor", "PhD", "PhD",
"PhD", "bachelor", "master", "highschool", "PhD", "highschool",
"bachelor", "master", "highschool", "highschool", "master",
"master", "bachelor", "PhD", "highschool", "PhD", "master",
"master", "master", "PhD", "highschool", "master", "master",
"highschool", "bachelor", "highschool", "bachelor", "PhD",
"bachelor", "highschool", "master", "highschool", "bachelor",
"bachelor"), region = c("North", "South", "North", "West",
"North", "West", "South", "West", "South", "West",
"South", "West", "East", "North", "West", "North", "North",
"West", "North", "East", "West", "South", "North", "North",
"East", "East", "North", "North", "West", "South", "West",
"West", "East", "West", "North", "West", "North", "East",
"North", "West", "South", "South", "East", "North", "West",
"West", "East", "North", "East"), marital_status = c("divorced",
"single", "divorced", "divorced", "divorced", "divorced",
"divorced", "married", "divorced", "married", "divorced",
"widowed", "married", "single", "widowed", "widowed", "single",
"divorced", "widowed", "widowed", "single", "married", "single",
"married", "widowed", "married", "single", "single", "widowed",
"married", "widowed", "divorced", "single", "married", "single",
"widowed", "widowed", "married", "widowed", "divorced", "married",
"married", "divorced", "single", "married", "widowed", "divorced",
"divorced", "single", "divorced)), row.names = c(NA, -50L
), class = c("tbl_df", "tbl", "data.frame"))

library(UtilsDataRSV)
see <- given_data |> view_cols()

```

```

[1] "id"
[1] "id_4"  "id_7"  "id_44" "id_48" "id_50" "id_30" "id_21" "id_31" "id_41"

```

```

[10] "id_6"  "id_42" "id_46" "id_25" "id_8"   "id_18" "id_22" "id_45" "id_33"
[19] "id_28" "id_12"
[1] "30 unique entries not displayed"
[1] "-----"
[1] "age"
[1] 69 76 31 24 64
[1] "-----"
[1] "gender"
[1] "male"   "female" "femal"
[1] "-----"
[1] "height"
[1] 150.5 161.0 165.4 174.4     NA
[1] "-----"
[1] "weight"
[1] 57.6 75.1 66.1 81.0 79.1
[1] "-----"
[1] "blood_type"
[1] "B"    "AB"   "A"    "O"
[1] "-----"
[1] "disease_status"
[1] "diseased" "Healthy" "healthy"
[1] "-----"
[1] "cholesterol"
[1] 160 213 162 195 230
[1] "-----"
[1] "glucose"
[1] 81 99 87 86 NA
[1] "-----"
[1] "smoker"
[1] "no"   "yes"
[1] "-----"
[1] "exercise"
[1] "regular"   "occasional" "none"
[1] "-----"
[1] "income"
[1] 25654 37940 88662 51178 56261
[1] "-----"
[1] "education"
[1] "master"      "bachelor"    "PhD"        "highschool"
[1] "-----"
[1] "region"
[1] "South" "North" "West"  "East"
[1] "-----"

```

```
[1] "marital_status"
[1] "single"    "widowed"   "married"   "divorced"
[1] "-----"
```

Warning: Not all unique entries displayed for these non-numeric cols: id

```
print('From the above we can see that the gender column has a typo: \'femal\' instead of \'fe
```

```
[1] "From the above we can see that the gender column has a typo: 'femal' instead of 'female'
```

```
print('The height column has some missing entries.')
```

```
[1] "The height column has some missing entries."
```

```
print('The disease_status column has a typo, \'Healthy\' and \'healthy\'.')
```

```
[1] "The disease_status column has a typo, 'Healthy' and 'healthy'."
```

```
print('The glucose column has some missing entries.')
```

```
[1] "The glucose column has some missing entries."
```