

Final Version of Assignment One

Amahle Nzimande

2026-02-09

Table of contents

Assignment One	3
Welcome	3
1 Project One	4
2 Day Three Practical Q3 & Q4	8
3 Day Four Practical	11

Assignment One

Welcome

Dear fellow survivor of this course

If you are reading this, we have both made it this far, congratulations to us! I hope this assignment makes sense, runs without errors, and causes minimal emotional damage.

May your marking be fair, your coffee strong, and your rubric generous.

1 Project One

QUESTION ONE

```
down <- nrow(airquality)
across <- ncol(airquality)
sum <- 0

print('The following rows have missing information: ')
```

```
[1] "The following rows have missing information: "
```

```
for(i in 1:down){
  for(j in 1:across){
    if(is.na(airquality[i,j])){
      sum <- sum + 1
      print(i)
      break}
  }
}
```

```
[1] 5
[1] 6
[1] 10
[1] 11
[1] 25
[1] 26
[1] 27
[1] 32
[1] 33
[1] 34
[1] 35
[1] 36
[1] 37
[1] 39
```

```
[1] 42  
[1] 43  
[1] 45  
[1] 46  
[1] 52  
[1] 53  
[1] 54  
[1] 55  
[1] 56  
[1] 57  
[1] 58  
[1] 59  
[1] 60  
[1] 61  
[1] 65  
[1] 72  
[1] 75  
[1] 83  
[1] 84  
[1] 96  
[1] 97  
[1] 98  
[1] 102  
[1] 103  
[1] 107  
[1] 115  
[1] 119  
[1] 150
```

```
print(paste0('In total there are ', sum, ' rows with missing information.' ))
```

```
[1] "In total there are 42 rows with missing information."
```

QUESTION TWO

```
my_table <- data.frame(  
  Column = c('Temperature', 'Ozone'),  
  Mean = c(mean(airquality[,4]), mean(airquality[,1], na.rm=TRUE)),  
  SD = c(sd(airquality[,4]), sd(airquality[,1], na.rm=TRUE)),  
  Min = c(min(airquality[,4]), min(airquality[,1], na.rm=TRUE)),  
  Max = c(max(airquality[,4]), max(airquality[,1], na.rm=TRUE))
```

```
)  
  
print(my_table)
```

	Column	Mean	SD	Min	Max
1	Temperature	77.88235	9.46527	56	97
2	Ozone	42.12931	32.98788	1	168

QUESTION THREE

```
data(cars)  
  
Y <- cars$dist  
X <- cbind(1, cars$speed)  
  
my_funct <- function(design, response) {  
  
  a <- t(design) %*% design  
  b <- solve(a)  
  c <- t(design) %*% response  
  d <- b %*% c  
  
  return(d)  
}  
  
print(my_funct(X, Y))
```

```
[,1]  
[1,] -17.579095  
[2,] 3.932409
```

QUESTION FOUR

```
model <- lm(cars$dist ~ cars$speed, data=cars )  
summary(model)
```

Call:
lm(formula = cars\$dist ~ cars\$speed, data = cars)

Residuals:

Min	1Q	Median	3Q	Max
-29.069	-9.525	-2.272	9.215	43.201

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-17.5791	6.7584	-2.601	0.0123 *
cars\$speed	3.9324	0.4155	9.464	1.49e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 15.38 on 48 degrees of freedom

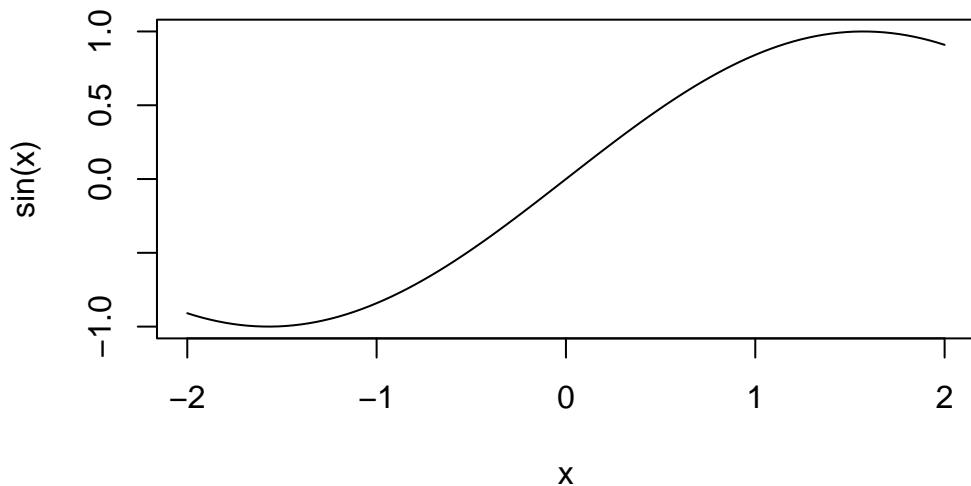
Multiple R-squared: 0.6511, Adjusted R-squared: 0.6438

F-statistic: 89.57 on 1 and 48 DF, p-value: 1.49e-12

2 Day Three Practical Q3 & Q4

QUESTION FOUR

```
curve(sin(x), from = -2, to = 2)
```



QUESTION FIVE

```
randoms <- rt(1000, 1)

#Creating a manual QQ-plot with 95% confidence interval

sorted_randoms <- sort(randoms)

n <- length(randoms)
i <- 1:n
```

```

# Calculate plotting positions (probabilities)
p <- (i - 3/8) / (n + 1/4)

# Generate normal quantiles by simulation (since we can't use qnorm)
large_normal_sample <- rnorm(1000000) # Large normal reference
theoretical_quantiles <- quantile(large_normal_sample, probs = p)

# Calculate 95% probability envelopes by simulation
n_sim <- 1000 # Number of simulations
envelope_matrix <- matrix(NA, nrow = n_sim, ncol = n)

# Simulate many normal samples of size n
for (j in 1:n_sim) {
  sim_sample <- rnorm(n)
  sim_sorted <- sort(sim_sample)
  envelope_matrix[j, ] <- sim_sorted
}

# Calculate 2.5% and 97.5% percentiles at each position
lower_envelope <- apply(envelope_matrix, 2, quantile, probs = 0.025)
upper_envelope <- apply(envelope_matrix, 2, quantile, probs = 0.975)

# Create the QQ-plot
par(mfrow = c(1, 2)) # Split plot window

# Manual QQ-plot
plot(theoretical_quantiles, sorted_randoms,
      xlab = "Theoretical Normal Quantiles",
      ylab = "Sample Quantiles",
      main = "Manual QQ-plot with 95% Envelopes",
      pch = 19, cex = 0.6,
      ylim = c(-5, 5), xlim = c(-4, 4),
      col = rgb(0, 0, 0, 0.7))

# Add reference line (y = x)
abline(a = 0, b = 1, col = "red", lwd = 2)

# Add 95% probability bands
lines(theoretical_quantiles, lower_envelope,
      col = "blue", lty = 2, lwd = 2)
lines(theoretical_quantiles, upper_envelope,
      col = "blue", lty = 2, lwd = 2)

```

```

# Add legend
legend("topleft",
       legend = c("Data", "Reference Line", "95% Envelope"),
       col = c("black", "red", "blue"),
       pch = c(19, NA, NA),
       lty = c(NA, 1, 2),
       lwd = c(NA, 2, 2))

# Check with car::qqPlot for comparison
library(car)

```


Loading required package: carData

```


qqPlot(randoms, envelope = 0.95, ylim = c(-5, 5),
       main = "car::qqPlot for Comparison",
       xlab = "Normal Quantiles",
       ylab = "Sample Quantiles")

```

Manual QQ-plot with 95% Envelope car::qqPlot for Comparison



Theoretical Normal Quantiles



Normal Quantiles

[1] 843 140

3 Day Four Practical

QUESTION ONE

```
library(tidyverse)
```

```
Warning: package 'tidyverse' was built under R version 4.3.3
```

```
Warning: package 'tidyr' was built under R version 4.3.3
```

```
Warning: package 'purrr' was built under R version 4.3.3
```

```
Warning: package 'dplyr' was built under R version 4.3.3
```

```
Warning: package 'forcats' was built under R version 4.3.3
```

```
Warning: package 'lubridate' was built under R version 4.3.3
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
v forcats   1.0.0     v stringr   1.6.0
v ggplot2   4.0.2     v tibble    3.2.1
v lubridate 1.9.4     v tidyr    1.3.1
v purrr    1.0.4
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()   masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to beco
```

```
library(nycflights13)
```

```
Warning: package 'nycflights13' was built under R version 4.3.3
```

```
#The line below renames the first 10 columns of flights and prints the flights data set with

flights |> rename(Year = year, Month = month, Day = day, Dep_Time = dep_time, Sched_Dep_Time

# A tibble: 336,776 x 19
   Year Month   Day Dep_Time Sched_Dep_Time Dep_Delay Arr_Time Sched_Arr_Time
   <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1  2013     1     2     2130        2130       0         1        18
2  2013     1    11     2157        2000      117        1      2208
3  2013     1    11     2253        2249       4         1      2357
4  2013     1    14     2122        2130      -8         1        2
5  2013     1    14     2246        2250      -4         1        7
6  2013     1    15     2304        2245      19         1      2357
7  2013     1    16     2018        2025      -7         1      2329
8  2013     1    16     2303        2245      18         1      2357
9  2013     1    19     2107        2110      -3         1      2355
10 2013     1    22     2246        2249      -3         1      2357
# i 336,766 more rows
# i 11 more variables: Arr_Delay <dbl>, Carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

QUESTION TWO

```
#Creates flight1 tibble that contains flights that occurred in month 1 ONLY
flight1 <- flights |> filter(month == 1)

#Gets the average distance of each carrier
carrier_dist_vec_mean <- flight1 |> group_by(carrier) |> summarise(MEAN = mean(distance, na.rm = TRUE))

#Gets the standard deviation of the distance for each carrier
carrier_dist_vec_sd <- flight1 |> group_by(carrier) |> summarise(SD = sd(distance, na.rm = TRUE))

#Creates one tibble with both the mean and sd distance for each carrier
dist_tbl <- carrier_dist_vec_mean |> left_join(carrier_dist_vec_sd)
```

Joining with `by = join_by(carrier)`

```
#Arranges the carrier rows from that with the least to the highest mean distance
dist_tbl |> arrange(MEAN)
```

```
# A tibble: 16 x 3
  carrier   MEAN     SD
  <chr>    <dbl>  <dbl>
1 YV        229     0
2 9E       476.  334.
3 EV        522.  294.
4 US        536.  553.
5 MQ        566.  223.
6 FL       691.  142.
7 00        733    NA
8 WN       942.  496.
9 B6      1062.  681.
10 DL     1220.  644.
11 AA     1350.  626.
12 UA     1462.  778.
13 F9      1620     0
14 AS      2402     0
15 VX     2495.  98.2
16 HA      4983     0
```

QUESTION THREE

```
#####
NA #####
carrier00 <- flight1 |> filter(carrier == '00') |> select(carrier, distance)

carrier00

# A tibble: 1 x 2
  carrier distance
  <chr>    <dbl>
1 00        733

print('Explanation of NA: Carrier 00 has NA as the standard deviation of the distance, that is')

[1] "Explanation of NA: Carrier 00 has NA as the standard deviation of the distance, that is"

#####
ZERO #####
carrierYV <- flight1 |> filter(carrier == 'YV') |> select(carrier, distance)

carrierYV
```

```

# A tibble: 46 x 2
  carrier distance
  <chr>     <dbl>
1 YV         229
2 YV         229
3 YV         229
4 YV         229
5 YV         229
6 YV         229
7 YV         229
8 YV         229
9 YV         229
10 YV        229
# i 36 more rows

print('Explanaiton of ZERO: The standard deviation of the distance for the YV carrier is zero')

[1] "Explanaiton of ZERO: The standard deviation of the distance for the YV carrier is zero"

```

QUESTION FOUR

```

#Getting the Average Departure delay for each carrier in each month in the long format
switched <- flights |> group_by(carrier, month) |> summarise(Av_Dep_Delay = mean(dep_delay, na.rm = TRUE))

`summarise()` has grouped output by 'carrier'. You can override using the
`.groups` argument.

#Converting into wide data
switched2 <- switched |> pivot_wider(names_from = carrier, values_from = Av_Dep_Delay)

switched2

# A tibble: 12 x 17
  month `9E`    AA    AS    B6    DL    EV    F9    FL    HA    MQ    OO
  <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1     1 16.9   6.93  7.35  9.49  3.85 24.2   10   1.97 54.4   6.49 67
2     2 16.5   8.28  0.722 13.8   5.54 21.5   29.8  5.18 17.4   8.09 NA
3     3 13.4   8.70  8.42  14.2   9.93 26.2   16.8  17.3  1.16  7.19 NA
4     4 13.6  11.7  11.3   15.2   8.17 22.8   24.6  13.1 -2.1   13.7 NA
5     5 22.7   9.66  6.77  9.78  9.74 20.2   35.9  19.2 -1.45 13.9 NA

```

```

6      6 29.0 14.6 13.1 20.4 18.7 25.5 29.4 38.8 1.47 20.8 61
7      7 31.4 12.1 2.42 24.9 20.6 26.5 31.8 41.2 -1.71 20.7 NA
8      8 17.3 7.17 2.87 15.7 9.85 16.3 22.2 23.4 1.68 10.1 64
9      9 7.75 5.69 -4.52 6.63 5.53 8.24 8.26 16.9 -5.44 5.35 -4.94
10     10 9.33 3.00 0.677 2.96 3.42 13.4 9.70 13.7 -5.10 4.48 NA
11     11 7.56 3.10 3.08 3.52 2.85 9.83 13.5 16.9 -5.44 3.28 0.8
12     12 19.8 11.7 18.0 17.0 10.8 27.9 13.1 26.1 -3.14 12.7 NA
# i 5 more variables: UA <dbl>, US <dbl>, VX <dbl>, WN <dbl>, YV <dbl>

```

QUESTION FIVE

```

#Getting number of all flights with a delayed departure time and arrived with NO delay
meet_Condition <- flights |> filter(dep_delay > 0) |> filter(arr_delay <= 0) |> count()

#Getting number of total flights
total_Flights <- flights |> count()

#Calculating the proportion
prop_Delayed <- meet_Condition/total_Flights

#Renaming the column
prop_Delayed <- prop_Delayed |> rename(Prop_Delayed = n)

prop_Delayed

```

```

Prop_Delayed
1    0.1052391

```

QUESTION SIX

```

routes <- flights |> group_by(origin, dest) |> summarize(n_airlines = n_distinct(carrier)) |>

`summarise()` has grouped output by 'origin'. You can override using the
`.groups` argument.

routes

# A tibble: 128 x 3
# Groups:   origin [3]
  origin dest  n_airlines
  <fct>  <fct> <dbl>
1 ATL    BNA    1.00 
2 ATL    CLT    1.00 
3 ATL    EWR    1.00 
4 ATL    IAH    1.00 
5 ATL    MIA    1.00 
6 ATL    PBI    1.00 
7 ATL    PDX    1.00 
8 ATL    SFO    1.00 
9 ATL    TPA    1.00 
10 ATL   VIE    1.00 
11 ATL   YYZ    1.00 
12 BNA   ATL    1.00 
13 BNA   CLT    1.00 
14 BNA   EWR    1.00 
15 BNA   IAH    1.00 
16 BNA   MIA    1.00 
17 BNA   PBI    1.00 
18 BNA   PDX    1.00 
19 BNA   SFO    1.00 
20 BNA   TPA    1.00 
21 BNA   VIE    1.00 
22 BNA   YYZ    1.00 
23 CLT   ATL    1.00 
24 CLT   BNA    1.00 
25 CLT   EWR    1.00 
26 CLT   IAH    1.00 
27 CLT   MIA    1.00 
28 CLT   PBI    1.00 
29 CLT   PDX    1.00 
30 CLT   SFO    1.00 
31 CLT   TPA    1.00 
32 CLT   VIE    1.00 
33 CLT   YYZ    1.00 
34 EWR   ATL    1.00 
35 EWR   BNA    1.00 
36 EWR   CLT    1.00 
37 EWR   IAH    1.00 
38 EWR   MIA    1.00 
39 EWR   PBI    1.00 
40 EWR   PDX    1.00 
41 EWR   SFO    1.00 
42 EWR   TPA    1.00 
43 EWR   VIE    1.00 
44 EWR   YYZ    1.00 
45 IAH   ATL    1.00 
46 IAH   BNA    1.00 
47 IAH   CLT    1.00 
48 IAH   EWR    1.00 
49 IAH   MIA    1.00 
50 IAH   PBI    1.00 
51 IAH   PDX    1.00 
52 IAH   SFO    1.00 
53 IAH   TPA    1.00 
54 IAH   VIE    1.00 
55 IAH   YYZ    1.00 
56 MIA   ATL    1.00 
57 MIA   BNA    1.00 
58 MIA   CLT    1.00 
59 MIA   EWR    1.00 
60 MIA   IAH    1.00 
61 MIA   PBI    1.00 
62 MIA   PDX    1.00 
63 MIA   SFO    1.00 
64 MIA   TPA    1.00 
65 MIA   VIE    1.00 
66 MIA   YYZ    1.00 
67 PBI   ATL    1.00 
68 PBI   BNA    1.00 
69 PBI   CLT    1.00 
70 PBI   EWR    1.00 
71 PBI   IAH    1.00 
72 PBI   MIA    1.00 
73 PBI   PBI    1.00 
74 PBI   PDX    1.00 
75 PBI   SFO    1.00 
76 PBI   TPA    1.00 
77 PBI   VIE    1.00 
78 PBI   YYZ    1.00 
79 PDX   ATL    1.00 
80 PDX   BNA    1.00 
81 PDX   CLT    1.00 
82 PDX   EWR    1.00 
83 PDX   IAH    1.00 
84 PDX   MIA    1.00 
85 PDX   PBI    1.00 
86 PDX   PDX    1.00 
87 PDX   SFO    1.00 
88 PDX   TPA    1.00 
89 PDX   VIE    1.00 
90 PDX   YYZ    1.00 
91 SFO   ATL    1.00 
92 SFO   BNA    1.00 
93 SFO   CLT    1.00 
94 SFO   EWR    1.00 
95 SFO   IAH    1.00 
96 SFO   MIA    1.00 
97 SFO   PBI    1.00 
98 SFO   PDX    1.00 
99 SFO   SFO    1.00 
100 SFO   TPA    1.00 
101 SFO   VIE    1.00 
102 SFO   YYZ    1.00 
103 TPA   ATL    1.00 
104 TPA   BNA    1.00 
105 TPA   CLT    1.00 
106 TPA   EWR    1.00 
107 TPA   IAH    1.00 
108 TPA   MIA    1.00 
109 TPA   PBI    1.00 
110 TPA   PDX    1.00 
111 TPA   SFO    1.00 
112 TPA   TPA    1.00 
113 TPA   VIE    1.00 
114 TPA   YYZ    1.00 
115 VIE   ATL    1.00 
116 VIE   BNA    1.00 
117 VIE   CLT    1.00 
118 VIE   EWR    1.00 
119 VIE   IAH    1.00 
120 VIE   MIA    1.00 
121 VIE   PBI    1.00 
122 VIE   PDX    1.00 
123 VIE   SFO    1.00 
124 VIE   TPA    1.00 
125 VIE   VIE    1.00 
126 VIE   YYZ    1.00 
127 YYZ   ATL    1.00 
128 YYZ   BNA    1.00 
129 YYZ   CLT    1.00 
130 YYZ   EWR    1.00 
131 YYZ   IAH    1.00 
132 YYZ   MIA    1.00 
133 YYZ   PBI    1.00 
134 YYZ   PDX    1.00 
135 YYZ   SFO    1.00 
136 YYZ   TPA    1.00 
137 YYZ   VIE    1.00 
138 YYZ   YYZ    1.00

```

```

<chr> <chr> <int>
1 EWR   ATL      4
2 EWR   AUS      2
3 EWR   BDL      2
4 EWR   BNA      2
5 EWR   BOS      3
6 EWR   BWI      2
7 EWR   CHS      2
8 EWR   CLE      2
9 EWR   CLT      3
10 EWR  CVG      2
# i 118 more rows

```

```
#Joining the routes with at least two airlines with rest of the flights data
joint <- routes |> left_join(flights, by = join_by(origin, dest))
```

```
#Grouping to get each route+carrier combo then calc it's average arrival delay
joint2 <- joint |> group_by(origin, dest, carrier) |> summarise(Av_Arr_Delay = mean(arr_delay))
```

``summarise()` has grouped output by 'origin', 'dest'. You can override using the ` `.groups` argument.`

```
joint2
```

```

# A tibble: 343 x 4
# Groups:   origin, dest [128]
  origin dest  carrier Av_Arr_Delay
  <chr>  <chr> <chr>     <dbl>
1 EWR    ATL    9E        -6.25
2 EWR    ATL    DL       10.00
3 EWR    ATL    EV        19.5
4 EWR    ATL    UA        10.5
5 EWR    AUS    UA        4.28
6 EWR    AUS    WN       -11.2
7 EWR    BDL    EV        6.78
8 EWR    BDL    UA        22.6
9 EWR    BNA    EV        17.7
10 EWR   BNA   WN       -2.13
# i 333 more rows

```

```
# Airline with the best average arrival delay

best <- joint2 |> group_by(origin, dest) |> summarise(best_av_arr_delay = min(Av_Arr_Delay))

`summarise()` has grouped output by 'origin'. You can override using the
`.groups` argument.

best2 <- best |> left_join(joint2, join_by(best_av_arr_delay == Av_Arr_Delay)) |> select(!or

Warning in left_join(best, joint2, join_by(best_av_arr_delay == Av_Arr_Delay)): Detected an
to-many relationship between `x` and `y`.
i Row 17 of `x` matches multiple rows in `y`.
i Row 41 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected, set `relationship =
"many-to-many"` to silence this warning.

best2
```

	origin.x	dest.x	best_av_arr_delay	carrier
	<chr>	<chr>	<dbl>	<chr>
1	EWR	ATL	-6.25	9E
2	EWR	AUS	-11.2	WN
3	EWR	BDL	6.78	EV
4	EWR	BNA	-2.13	WN
5	EWR	BOS	-4.01	EV
6	EWR	BWI	5.95	WN
7	EWR	CHS	-14	UA
8	EWR	CLE	-3.71	EV
9	EWR	CLT	0.920	US
10	EWR	CVG	1.40	9E
	# i 124 more rows			

```
#Airline with the worst arrival delay

worst <- joint2 |> group_by(origin, dest) |> summarise(worst_av_arr_delay = max(Av_Arr_Delay))

`summarise()` has grouped output by 'origin'. You can override using the
`.groups` argument.
```

```
worst2 <- worst |> left_join(joint2, join_by(worst_av_arr_delay == Av_Arr_Delay)) |> select(
```

Warning in left_join(worst, joint2, join_by(worst_av_arr_delay == Av_Arr_Delay)): Detected an
to-many relationship between `x` and `y`.
i Row 75 of `x` matches multiple rows in `y`.
i Row 194 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected, set `relationship =
"many-to-many"` to silence this warning.

```
worst2
```

```
# A tibble: 130 x 4  
  origin.x dest.x worst_av_arr_delay carrier  
  <chr>     <chr>          <dbl> <chr>  
1 EWR       ATL            19.5  EV  
2 EWR       AUS            4.28  UA  
3 EWR       BDL            22.6  UA  
4 EWR       BNA            17.7  EV  
5 EWR       BOS            6.87  B6  
6 EWR       BWI            20.1  EV  
7 EWR       CHS            16.2  EV  
8 EWR       CLE            5.97  UA  
9 EWR       CLT            20.5  EV  
10 EWR      CVG            21.2  EV  
# i 120 more rows
```

```
#Combining
```

```
results <- best2 |> left_join(worst2, join_by(origin.x, dest.x))
```

Warning in left_join(best2, worst2, join_by(origin.x, dest.x)): Detected an unexpected many-
to-many relationship between `x` and `y`.
i Row 77 of `x` matches multiple rows in `y`.
i Row 17 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected, set `relationship =
"many-to-many"` to silence this warning.

```
results
```

```

# A tibble: 136 x 6
  origin.x dest.x best_av_arr_delay carrier.x worst_av_arr_delay carrier.y
  <chr>     <chr>           <dbl> <chr>           <dbl> <chr>
1 EWR       ATL            -6.25  9E              19.5  EV
2 EWR       AUS            -11.2   WN             4.28  UA
3 EWR       BDL             6.78  EV              22.6  UA
4 EWR       BNA            -2.13   WN             17.7  EV
5 EWR       BOS             -4.01  EV              6.87  B6
6 EWR       BWI             5.95   WN             20.1  EV
7 EWR       CHS            -14    UA              16.2  EV
8 EWR       CLE            -3.71   EV              5.97  UA
9 EWR       CLT             0.920  US             20.5  EV
10 EWR      CVG            1.40   9E             21.2  EV
# i 126 more rows

```

```

results2 <- results |> mutate(diff = sqrt((best_av_arr_delay - worst_av_arr_delay)**2))

results3 <- results2 |> summarise(greatest_diff = max(diff, na.rm = TRUE)) |> left_join(resu

results3

```

```

# A tibble: 1 x 7
  greatest_diff origin.x dest.x best_av_arr_delay carrier.x worst_av_arr_delay
  <dbl> <chr>     <chr>           <dbl> <chr>           <dbl>
1        127.  JFK      ATL            1.40  9E          128
# i 1 more variable: carrier.y <chr>

```

```

#Trying to find the reason

reason <- results |> mutate(diff = sqrt((best_av_arr_delay - worst_av_arr_delay)**2))

reason2 <- reason |> summarise(terrible = max(worst_av_arr_delay, na.rm = TRUE))

reason2

```

```

# A tibble: 1 x 1
  terrible
  <dbl>
1      128

```

```
print('The above code shows that the worst average arrival delay was 128, it so happens that t
```

```
[1] "The above code shows that the worst average arrival delay was 128, it so happens that t
```

QUESTION SEVEN

```
given_data <- structure(list(id = c("id_1", "id_2", "id_3", "id_4", "id_5",
"id_6", "id_7", "id_8", "id_9", "id_10", "id_11", "id_12", "id_13",
"id_14", "id_15", "id_16", "id_17", "id_18", "id_19", "id_20",
"id_21", "id_22", "id_23", "id_24", "id_25", "id_26", "id_27",
"id_28", "id_29", "id_30", "id_31", "id_32", "id_33", "id_34",
"id_35", "id_36", "id_37", "id_38", "id_39", "id_40", "id_41",
"id_42", "id_43", "id_44", "id_45", "id_46", "id_47", "id_48",
"id_49", "id_50"), age = c(50L, 34L, 70L, 33L, 22L, 61L, 69L,
73L, 62L, 56L, 71L, 33L, 73L, 44L, 45L, 46L, 24L, 70L, 46L, 76L,
47L, 76L, 28L, 48L, 54L, 27L, 45L, 26L, 61L, 28L, 38L, 55L, 33L,
36L, 62L, 58L, 72L, 31L, 34L, 51L, 61L, 64L, 26L, 28L, 60L, 29L,
42L, 46L, 79L, 72L), gender = c("male", "male", "male", "female",
"female", "male", "female", "male", "male", "female", "female",
"male", "male", "female", "male", "male", "male", "female",
"male", "male", "male", "female", "femal", "male", "female",
"female", "female", "male", "female", "female", "female",
"male", "female", "male", "female", "female", "male",
"female", "female", "male", "female", "male", "male",
"male", "female"), height = c(174.4, 197.7, 174.1, 194.5, NA,
180.4, 170.5, 157.4, 196.8, 165.1, 153, 197.4, 186, 157.1, 177.5,
197.7, 179.3, 170.2, 182.4, NA, 165.4, 161, 168.5, 199.2, 157.7,
154.6, 157.1, 184.5, 181, 194.6, 183.6, 186.9, 176.1, 183, 191.1,
189.3, 199, 172, 165.6, 170.5, 150.5, 159.2, 192.1, 161.6, 162,
153.8, 162.3, 186.6, 192.4, 174.9), weight = c(69.4, 62.3, 55.6,
69.5, 78.6, 60.8, 72.2, 60.9, 75.1, 67.7, 82.5, 68.7, 67.8, 76.7,
87, 61.1, 70.6, 63.3, 81.5, 59.2, 93.2, 87.3, 83.4, 80.9, 68.6,
76.5, 93.7, 79.1, 92, 65.6, 85.4, 63.3, 79.7, 74.1, 63.3, 78.2,
95.7, 95.1, 63.7, 66.1, 99.3, 81, 96.9, 73.3, 70.3, 83, 57.6,
78.6, 61.9, 98.1), blood_type = c("O", "A", "O", "O", "B", "AB",
"O", "O", "O", "AB", "A", "O", "O", "B", "A", "B", "AB",
"O", "AB", "A", "AB", "O", "B", "A", "A", "B", "AB", "A", "B",
"B", "A", "O", "O", "B", "B", "O", "A", "A", "B", "A", "O", "AB",
"A", "A", "O", "O", "B", "A", "O"), disease_status = c("diseased",
"healthy", "healthy", "healthy", "healthy", "healthy", "diseased",
"healthy", "diseased", "Healthy", "diseased", "healthy", "diseased",
```

```
healthy", "diseased", "healthy", "healthy", "healthy", "healthy", "healthy",  
"healthy", "healthy", "diseased", "healthy", "diseased", "healthy",  
"healthy", "healthy", "diseased", "diseased", "diseased", "healthy",  
"diseased", "healthy", "healthy", "healthy", "healthy", "healthy",  
"diseased", "diseased", "diseased", "healthy", "healthy", "diseased",  
"diseased"), cholesterol = c(228, 223, 213, 198, 166, 151, 195,  
199, 189, 196, 221, 156, 185, 230, 234, 174, 185, 236, 235, 180,  
165, 220, 160, 153, 250, 153, 184, 242, 212, 179, 224, 233, 181,  
199, 220, 214, 214, 248, 191, 162, 203, 173, 199, 187, 248, 189,  
173, 212, 164, 247), glucose = c(96, 78, 101, 119, 103, 91, 86,  
NA, 77, 80, 115, 85, 88, 109, NA, 71, 90, 94, 91, 87, 113, 93,  
97, 118, 109, 80, 85, 119, 99, 108, 89, 108, 97, 116, 79, 84,  
75, 81, 119, NA, 106, 109, 75, 82, 84, 75, 76, 120, 119, 77),  
smoker = c("yes", "yes", "yes", "yes", "no", "yes", "no",  
"yes", "no", "no", "no", "no", "yes", "no", "yes",  
"yes", "yes", "yes", "yes", "yes", "yes", "yes", "yes", "no",  
"no", "yes", "yes", "no", "no", "yes", "no", "yes",  
"no", "yes", "no", "yes", "yes", "yes", "no", "no", "yes",  
"no", "no", "no", "no", "no", "yes"), exercise = c("occasional",  
"regular", "occasional", "regular", "none", "occasional",  
"regular", "none", "occasional", "none", "occasional", "none",  
"none", "regular", "occasional", "none", "regular", "regular",  
"none", "occasional", "none", "occasional", "occasional",  
"occasional", "regular", "occasional", "regular", "regular",  
"regular", "occasional", "occasional", "none", "none", "regular",  
"occasional", "occasional", "none", "none", "none",  
"occasional", "regular", "regular", "none", "regular", "occasional",  
"occasional", "none", "occasional", "regular"), income = c(84820L,  
81547L, 22588L, 72490L, 74533L, 25338L, 41469L, 57315L, 63629L,  
88662L, 62615L, 56261L, 58499L, 82232L, 77584L, 77275L, 38468L,  
54510L, 91326L, 78611L, 31402L, 29586L, 21441L, 58269L, 84173L,  
88295L, 37940L, 43750L, 69750L, 92356L, 82518L, 91455L, 68866L,  
51178L, 68275L, 27689L, 35418L, 81318L, 62405L, 86851L, 25654L,  
47553L, 74474L, 51409L, 22607L, 55360L, 96351L, 21516L, 41927L,  
55810L), education = c("master", "bachelor", "PhD", "master",  
"bachelor", "highschool", "PhD", "highschool", "PhD", "PhD",  
"bachelor", "highschool", "master", "bachelor", "PhD", "PhD",  
"PhD", "bachelor", "master", "highschool", "PhD", "highschool",  
"bachelor", "master", "highschool", "highschool", "master",  
"master", "bachelor", "PhD", "highschool", "PhD", "master",  
"master", "master", "PhD", "highschool", "master", "master",
```

```

"highschool", "bachelor", "highschool", "bachelor", "PhD",
"bachelor", "highschool", "master", "highschool", "bachelor",
"bachelor"), region = c("North", "South", "North", "West",
"North", "West", "South", "South", "West", "South", "West",
"South", "West", "East", "North", "West", "North", "North",
"West", "North", "East", "West", "South", "North", "North",
"East", "East", "North", "North", "West", "South", "West",
"West", "East", "West", "North", "West", "North", "East",
"North", "West", "South", "South", "East", "North", "West",
"West", "East", "North", "East"), marital_status = c("divorced",
"single", "divorced", "divorced", "divorced", "divorced",
"divorced", "married", "divorced", "married", "divorced",
"widowed", "married", "single", "widowed", "widowed", "single",
"divorced", "widowed", "widowed", "single", "married", "single",
"married", "widowed", "married", "single", "single", "widowed",
"married", "widowed", "divorced", "single", "married", "single",
"widowed", "widowed", "married", "widowed", "divorced", "married",
"married", "divorced", "single", "married", "widowed", "divorced",
"divorced", "single", "divorced")), row.names = c(NA, -50L
), class = c("tbl_df", "tbl", "data.frame"))

library(UtilsDataRSV)
see <- given_data |> view_cols()

```

```

[1] "id"
[1] "id_11" "id_17" "id_28" "id_48" "id_9"  "id_12" "id_44" "id_2"   "id_24"
[10] "id_32" "id_27" "id_35" "id_22" "id_43" "id_36" "id_3"   "id_49" "id_37"
[19] "id_41" "id_20"
[1] "30 unique entries not displayed"
[1] "-----"
[1] "age"
[1] 54 28 34 44 79
[1] "-----"
[1] "gender"
[1] "female" "male"   "femal"
[1] "-----"
[1] "height"
[1] 186.9 165.4 170.2 192.1    NA
[1] "-----"
[1] "weight"
[1] 78.6 93.7 96.9 79.1 61.1
[1] "-----"

```

```

[1] "blood_type"
[1] "O"   "B"   "A"   "AB"
[1] "-----"
[1] "disease_status"
[1] "diseased" "healthy"  "Healthy"
[1] "-----"
[1] "cholesterol"
[1] 156 195 166 180 165
[1] "-----"
[1] "glucose"
[1] 120 89 116 90 NA
[1] "-----"
[1] "smoker"
[1] "yes" "no"
[1] "-----"
[1] "exercise"
[1] "regular"    "occasional" "none"
[1] "-----"
[1] "income"
[1] 41469 47553 58269 38468 82232
[1] "-----"
[1] "education"
[1] "highschool" "master"      "PhD"        "bachelor"
[1] "-----"
[1] "region"
[1] "North" "South" "West"  "East"
[1] "-----"
[1] "marital_status"
[1] "divorced" "single"   "married"  "widowed"
[1] "-----"

```

Warning: Not all unique entries displayed for these non-numeric cols: id

```
print('LIST OF ERRORS: ')
```

```
[1] "LIST OF ERRORS: "
```

```
print('From the above we can see that the gender column has a typo: \'femal\' instead of \'femal\'')
```

```
[1] "From the above we can see that the gender column has a typo: 'femal' instead of 'female'"
```

```
print('The height column has some missing entries.')

[1] "The height column has some missing entries."

print('The disease_status column has a typo, \'Healthy\' and \'healthy\'.')
[1] "The disease_status column has a typo, 'Healthy' and 'healthy'."  
print('The glucose column has some missing entries.')
[1] "The glucose column has some missing entries."
```